

Macroeconometrics and Machine Learning : Lasso method

Kilian ANDRU, Gaëlle DORMION

Import data

```
library(MTS)
```

```
## Warning: le package 'MTS' a été compilé avec la version R 4.3.2
```

```
library(fbi)
library(glmnet)
```

```
## Warning: le package 'glmnet' a été compilé avec la version R 4.3.2
```

```
## Le chargement a nécessité le package : Matrix
```

```
## Loaded glmnet 4.1-8
```

```
# Setting start and end date of the analysis
start_date = "1960-01-01"
end_date = "2019-12-01" # We might want to leave out the covid period

# Load the most recent fred-md dataset

# Load transformed data
data = fredmd("https://files.stlouisfed.org/files/htdocs/fred-md/monthly/current.csv",
              date_start = as.Date(start_date), date_end = as.Date(end_date))

var_desc = fredmd_description

# Period
dates <- data$date
```

Data formatting

We remove useless columns and columns with NA values

```
#Remove the first column "date"
raw_data <- data[,seq(2,length(data))]
```

```
series <- names(data)

#Remove columns with NA values
names(raw_data[, colSums(is.na(raw_data)) > 0])

## [1] "ACOGNO"      "ANDENOx"      "TWEXAFEGSMTHx" "UMCSENTx"
## [5] "VIXCLSx"
```


```
x <- raw_data[, colSums(is.na(raw_data)) == 0]
```

```
#We check if there are still missing values
table(sapply(x, function(x) sum(length(which(is.na(x))))))
```

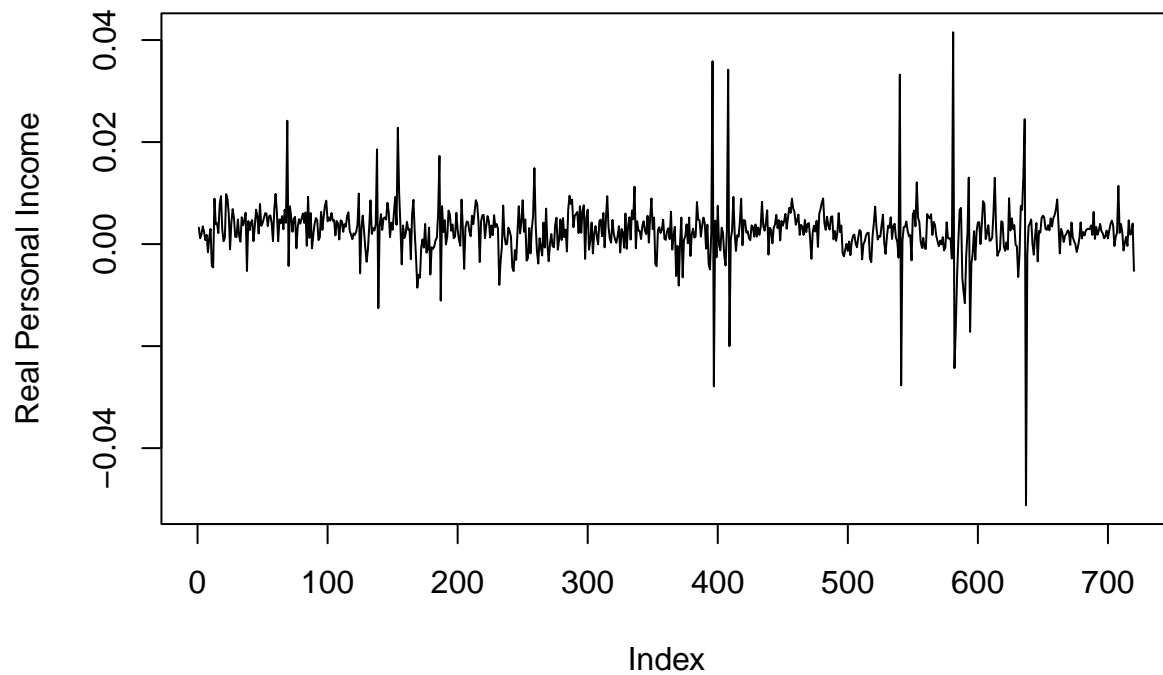
```
##
## 0
## 122
```

We plot the two time series of interest

```
plot(x$DPCERA3M086SBEA,type = "l",ylab = var_desc$description[var_desc$fred == "DPCERA3M086SBEA"])
```



```
plot(x$RPI, type = "l", ylab = var_desc$description[var_desc$fred == "RPI"])
```



Let standardize the data

```
x_standard <- x

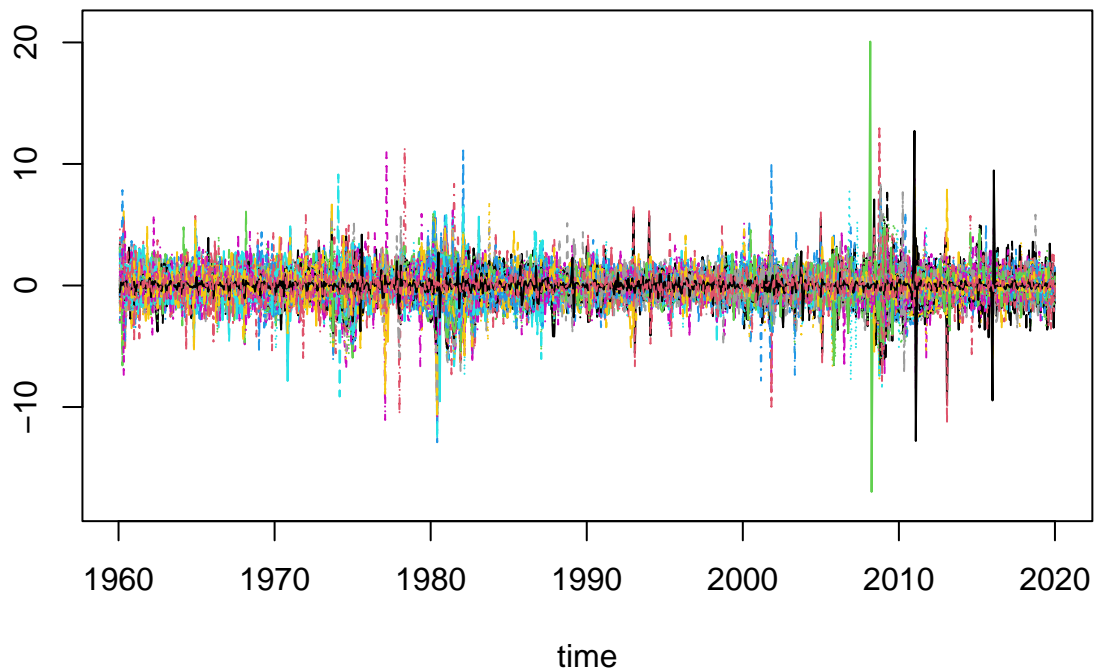
for (col in names(x)){
  mean <- mean(x[,col])
  st_d <- sd(x[,col])
  x_standard[,col] <- (x[,col] - mean) / st_d
}

#Time variable
time <- c(1:nrow(x_standard))/12 + 1960

#Number of period
nb_T <- nrow(x_standard)

#Number of variable
nb_N <- ncol(x_standard)

#Plot the series
MTSplot(x_standard,time)
```



The standardized data looks similar as what we got in class

Lasso regression on real personal consumption

function for lasso regression

```
Lasso_function <- function(p, r, alpha, variable, plot,residuals){
  if (p <= (r)){
    print("Must have p > r")
  }

  #outcome
  Y <- x_standard[(p+1):nb_T,which(names(x_standard)==variable)]
  #initialize endogenous variable
  X_st <- x_standard[p:(nb_T-1),which(names(x_standard)==variable)]
  if (p >= 2){
    for (i in seq(1,p-1)){
      Y_lag <- x_standard[seq(p-i,nb_T-1-i),which(names(x_standard)==variable)]
      X_st <- cbind(X_st,Y_lag)
    }
  }
  #exogenous variables and lag
  for (i in seq(0,r)){
    temp <- x_standard[(p-i):(nb_T-i-1),-which(names(x_standard)==variable)]
    X_st <- cbind(X_st,temp)
  }
}
```

```

X_st <- as.matrix(X_st)

#We use a cross validation to get the optimal lambda
cv_fit <- cv.glmnet(X_st,Y,alpha=alpha)
lambda <- cv_fit$lambda.min

#To get the residuals
if (residuals){
  resid <- Y - predict(cv_fit, newx = X_st)
  return(resid)
}

#We get the estimated lasso model
lasso_fit <- glmnet(X_st,Y,alpha=alpha,lambda = lambda)
#Plot if needed
if (plot){
  plot(glmnet(X_st,Y,alpha=alpha))
  plot(cv_fit)
}

return(lasso_fit)
}

```

We can now extract the coefficients to try to forecast the series, we start by getting the coefficients and the corresponding lag

```

list_coeff <- function(Lasso_fit){
  coefficients <- Lasso_fit$beta

  #For Y variable
  coef_Y <- coefficients[1:13]
  lag <- (1:13)
  coefficients_Y <- data.frame(coef_Y,lag)
  #Remove variables with null coeff
  coefficients_Y <- coefficients_Y[coefficients_Y$coef_Y !=0,]

  #For exogenous variables
  coef_X <- coefficients[seq(14,length(coefficients))]
  FRED <- row.names(Lasso_fit$beta)[seq(14,length(coefficients))]

  #We initialize the dataframe, there is 131 variables for each lag
  coefficients_X <- data.frame(coef_X = coef_X[seq(1,131)])
  coefficients_X$lag <- 1
  coefficients_X$FRED <- FRED[seq(1,131)]
  #Remove variables with null coeff
  coefficients_X <- coefficients_X[coefficients_X$coef_X !=0,]

  #For other lag
  for (i in (1:11)){
    coef_X <- coef_X[seq(i*131+1,(i+1)*131)]
    temp <- data.frame(coef_X)
    temp$lag <- i+1
  }
}

```

```

temp$FRED <- FRED[seq(i*131+1,(i+1)*131)]

#Remove variables with null coeff
temp <- temp[temp$coeff_X != 0,]

#Add to dataframe
coefficients_X <- rbind(coefficients_X, temp)
}

return (list(COEF_Y = coefficients_Y, COEF_X = coefficients_X))
}

```

Forecasting

Here we compute the standardized forecasts

```

forecasting <- function(Lasso_fit, variable, h, standard){

  #outcome
  Y <- x_standard[,which(names(x_standard)==variable)]
  #other variables
  X_st <- x_standard[,~which(names(x_standard)==variable)]

  list_coef <- list_coef(Lasso_fit)

  #We initialize the list of forecast values
  forecasts <- c()
  #for each period
  for (i in seq(1,h)){
    value <- 0

    #for the endogenous variables, each lag
    for (k in seq(1,nrow(list_coef$COEF_Y))){
      lag <- as.numeric(list_coef$COEF_Y$lag[k])
      #coeff
      coeff <- as.numeric(list_coef$COEF_Y$coeff_Y[k])

      #if we reuse a forecasted value
      if (i > lag){
        value <- value + forecasts[length(forecasts)-lag+1]*coeff
      }else{ #When we use value from the data set
        value <- value + Y[nb_T+ i -lag]*coeff
      }
    }
  }

  #For endogenous variables
  for (k in seq(1,nrow(list_coef$COEF_X))){
    lag <-list_coef$COEF_X$lag[k]
    #coeff
    coeff <- as.numeric(list_coef$COEF_X$coeff_X[k])
  }
}

```

```

    index <- which(names(X_st)== list_coef$COEF_X$FRED[k])

    #There is no forecasted value for exogenous, we only use the value it is in the dataset
    if (i <= lag){
      value <- value + X_st[nb_T+ i -lag, index]*coeff
    }
  }
  forecasts <- c(forecasts,value)
}
#If we want standardized data
if (standard) {
  return(forecasts)
}else{
  my <- mean(x[,which(names(x)==variable)])
  sy <- sd(x[,which(names(x)==variable)])
  return(forecasts*sy + my)
}
}

```

Forecasting error using naive forecasting

```

forecasting_error <- function(Lasso_fit, variable, h,sigma){
  Forecasts <- forecasting(Lasso_fit, variable, h, standard = FALSE)

  min <- c()
  max <- c()
  for (i in seq(1,h)){
    error <- sd(x[,which(names(x)==variable)])*sigma*(i)^(1/2)
    min <- c(min,Forecasts[1] - error)
    max <- c(max,Forecasts[1] + error)
  }

  return(data.frame(Forecasts,min,max))
}

```

Plot the forecasted values

```

library(ggplot2)
plot_forecast <- function(forecasts_data, variable, start){

  Forecasts <- x[,which(names(x)==variable)]

  #forecast horizon
  h <- nrow(forecasts_data)

  #time variable
  time <- c(1:(length(Forecasts)+h))/12 + 1960

  #Data to plot
  data_plot <- data.frame(Forecasts)
  data_plot$min <- NA
  data_plot$max <- NA

```

```

#Adding forecast
data_plot <- rbind(data_plot,forecasts_data)

data_plot$time <- time

data_plot <- data_plot[data_plot$time >= start,]

p <- ggplot(data=data_plot, aes(time, Forecasts)) +
  geom_ribbon(aes(ymin = min, ymax = max), fill = "steelblue2") +
  geom_line(color = "firebrick", size = 0.1)

return(p)
}

```

Application to variables

Real personal consumption

Now we try to find the best Lasso-VAR model to fit real personal consumption.


We take p and r large, useless variables will be removed with the lasso.


```

p = 13
r =12
alpha = 1
variable = "DPCERA3M086SBEA"

Lasso_fit <- Lasso_function(p = p, r = r, alpha = alpha, variable = variable, plot = TRUE, residuals = 1)

```



```
print(Lasso_fit$lambda)
```

```
## [1] 0.04243689
```

```
print((log(nb_N-1)/nb_T)^(1/2))
```

```
## [1] 0.08161385
```

```
print(Lasso_fit$dev.ratio)
```

```
## [1] 0.3972482
```

We extract the residuals to compute the forecasting error

```
residual <- Lasso_function(p = p, r = r, alpha = alpha, variable = variable, plot = FALSE, residuals = 'residuals')
sigma <- sd(residual)
```

Number of variables kept, for the outcome

```
coefficients_Y <- Lasso_fit$beta[1:13]
length(coefficients_Y[coefficients_Y != 0])
```

```
## [1] 3
```

Number of variables kept, for the exogenous variables

```
coefficients_X <- Lasso_fit$beta[seq(14,length(Lasso_fit$beta))]  
length(coefficients_X[coefficients_X !=0])
```

```
## [1] 98
```

We get the variables with higher coefficient in absolute value (X_st correspond to Y_lag of order 1)


```
VAR_names <- row.names(Lasso_fit$beta)  
coeffs <- Lasso_fit$beta[seq(1,length(Lasso_fit$beta))]  
table_coefficients <- data.frame(VAR_names, coeffs)  
  
lags <- c(1,2,3,4,5,6,7,8,9,10,11,12,13)  
for (i in seq(1,13)){  
  for (k in seq(1,121)){  
    lags <- c(lags,i)  
  }  
}  
  
table_coefficients$lags <- lags  
table_coefficients$abs_coff <- abs(table_coefficients$coeffs)  
  
table_coefficients <-table_coefficients[order(table_coefficients$abs_coff,decreasing = TRUE),]  
head(table_coefficients)
```

```
##      VAR_names      coeffs lags  abs_coff  
## 1      X_st -0.20597404    1 0.20597404  
## 572  FEDFUNDS -0.08218649    5 0.08218649  
## 76    M2REAL  0.07517290    1 0.07517290  
## 869  IPDCONGD  0.07195246    8 0.07195246  
## 1122 HWIURATIO 0.06957909   10 0.06957909  
## 217      BAA -0.06143196    2 0.06143196
```

We plot the forecasted values


```
forecast_data <- forecasting_error(Lasso_fit, variable, h= 15,sigma)  
plot_forecast(forecast_data, variable, start=1960)
```

```
## Warning: Using 'size' aesthetic for lines was deprecated in ggplot2 3.4.0.  
## i Please use 'linewidth' instead.  
## This warning is displayed once every 8 hours.  
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was  
## generated.
```



We focus on a smaller period

```
plot_forecast(forecast_data, variable, start=2010)
```




Real personal income


Now we try on real personal income

We take p and r large, useless variables will be removed with the lasso

```
p = 13  
r = 12  
alpha = 1  
variable = "RPI"
```

```
Lasso_fit <- Lasso_function(p = p, r = r, alpha = alpha, variable = variable, plot = TRUE, residuals = 1)
```





```
print(Lasso_fit$lambda)
```

```
## [1] 0.1464438
```

```
print((log(nb_N-1)/nb_T)^(1/2))
```

```
## [1] 0.08161385
```

```
print(Lasso_fit$dev.ratio)
```

```
## [1] 0.1038434
```

We extract the residuals to compute the forecasting error

```
residual <- Lasso_function(p = p, r = r, alpha = alpha, variable = variable, plot = FALSE, residuals = 'residuals')
sigma <- sd(residual)
```

Number of variables kept, for the outcome

```
coefficients_Y <- Lasso_fit$beta[1:13]
length(coefficients_Y[coefficients_Y != 0])
```

```
## [1] 1
```

Number of variables kept, for the exogenous variables

```
coefficients_X <- Lasso_fit$beta[seq(14,length(Lasso_fit$beta))]  
length(coefficients_X[coefficients_X !=0])
```

```
## [1] 4
```


We get the variables with higher coefficient in absolute value (X_st correspond to Y_lag of order 1)

```
VAR_names <- row.names(Lasso_fit$beta)  
coeffs <- Lasso_fit$beta[seq(1,length(Lasso_fit$beta))]  
table_coefficients <- data.frame(VAR_names, coeffs)  
  
lags <- c(1,2,3,4,5,6,7,8,9,10,11,12,13)  
for (i in seq(1,13)){  
  for (k in seq(1,121)){  
    lags <- c(lags,i)  
  }  
}  
  
table_coefficients$lags <- lags  
table_coefficients$abs_coff <- abs(table_coefficients$coeffs)  
  
table_coefficients <-table_coefficients[order(table_coefficients$abs_coff,decreasing = TRUE),]  
head(table_coefficients)
```

```
##      VAR_names      coeffs lags  abs_coff  
## 321  NONBORRES 0.17551450   3 0.17551450  
##  48    MANEMP 0.02610497   1 0.02610497  
## 1107   INDPRO 0.01546044  10 0.01546044  
##  50   NDMANEMP 0.01420956   1 0.01420956  
##  12      Y_lag 0.01151505  12 0.01151505  
##   1       X_st 0.00000000   1 0.00000000
```

We plot the forecasted values

```
forecast_data <- forecasting_error(Lasso_fit, variable, h= 15,sigma)  
plot_forecast(forecast_data, variable, start=1960)
```

We focus on a smaller period

```
plot_forecast(forecast_data, variable, start=2010)
```

